


# Machine Learning Model Workflow Report



**Report Prepared by:**

**Aya Hamdy**

# Agenda

## 1. Introduction

- Overview of the Machine Learning model workflow
- Objective of the report

## 2. Data Preparation

- Data pre-processing
- Handling outliers
- Data augmentation techniques
- Standardizing Data

## 3. Model Selection with PyCaret

- Overview of PyCaret
- Rationale for using PyCaret
- Algorithm comparison

## 4. Modelling

- Gradient Boosting Classifier (GBC)
- Logistic Regression
- Support Vector Classifier (SVC)

## 5. TensorBoard Workflow

- Visualization of model metrics
- Workflow tracking

## 6. Conclusion

- Summary of findings

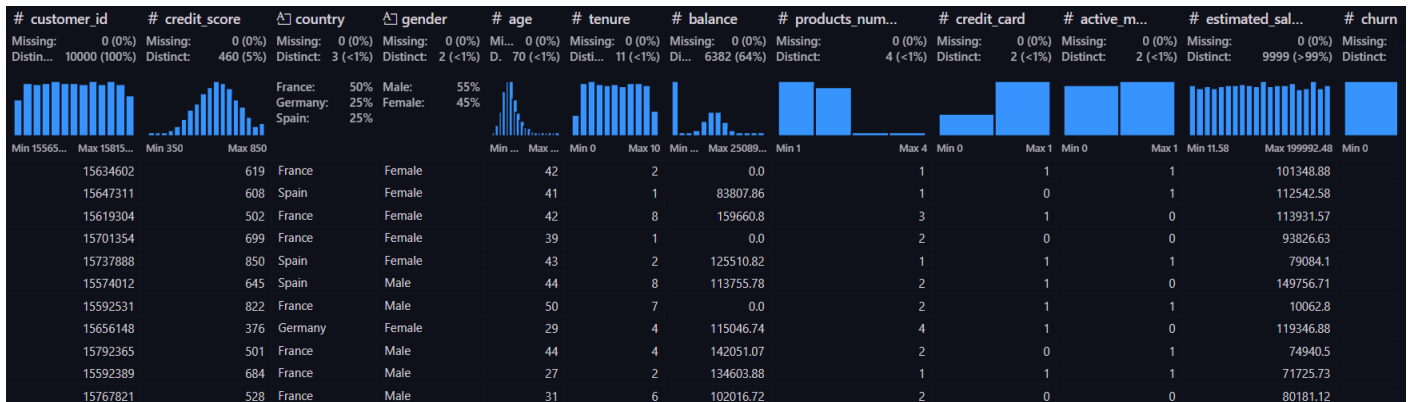
# 1. Introduction

This report outlines the workflow followed to build a Machine Learning model, covering data preparation, model selection, and evaluation. The primary objective was to develop a robust model that accurately classifies data using techniques like Logistic Regression, Support Vector Classifier (SVC), and Gradient Boosting Classifier. PyCaret was employed for algorithm comparison, and TensorBoard was used to track the workflow visually.

## 2. Data Preparation

### 2.1. Data Pre-Processing

The dataset consists of 10,000 entries and 11 columns, including features like `credit_score`, `age`, `balance`, and the target variable `churn`. The data was cleaned to remove missing values and inconsistencies.



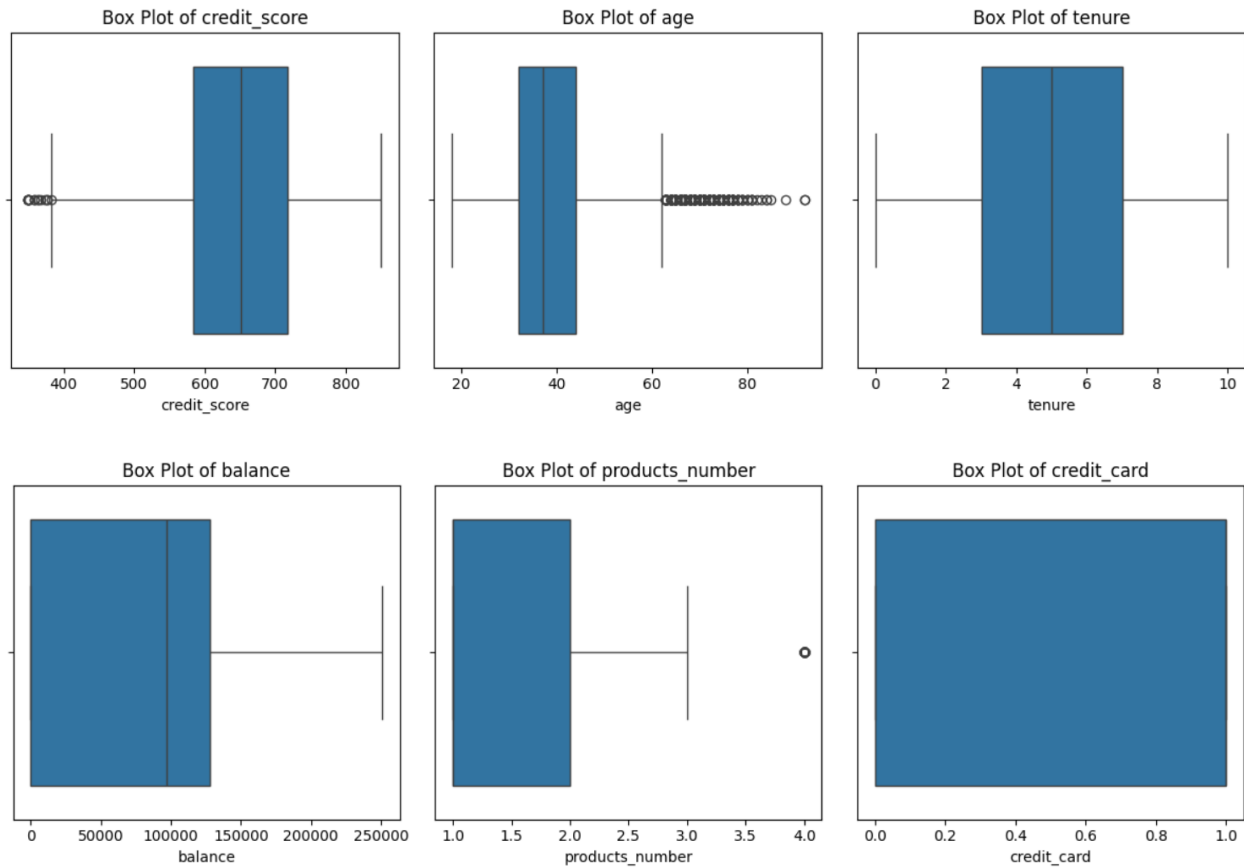
### Data Information:

```
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   credit_score     10000 non-null   int64
1   country          10000 non-null   object
2   gender           10000 non-null   object
3   age              10000 non-null   int64
4   tenure           10000 non-null   int64
5   balance          10000 non-null   float64
6   products_number  10000 non-null   int64
7   credit_card      10000 non-null   int64
8   active_member    10000 non-null   int64
9   estimated_salary  10000 non-null   float64
10  churn            10000 non-null   int64
dtypes: float64(2), int64(7), object(2)
```

## 2.2. Handling Outliers

Outliers were identified in the dataset as follows:

```
Data Shape: (10000, 11)
Total Outliers in credit_score: 15 -- 0.15%
Total Outliers in age: 359 -- 3.59%
Total Outliers in tenure: 0 -- 0.0%
Total Outliers in balance: 0 -- 0.0%
Total Outliers in products_number: 60 -- 0.6%
Total Outliers in credit_card: 0 -- 0.0%
Total Outliers in active_member: 0 -- 0.0%
Total Outliers in estimated_salary: 0 -- 0.0%
Total Outliers in churn: 2037 -- 20.37%
```





## Some Observations from above plots:

- We can see that data is highly imbalanced. Almost **80%** of our data is from class 0 (not exited) and **20%** data is from class 1 (exited).

```
1 # Check the class distribution before augmentation
2 print("class distribution before augmentation: ")
3 print(data['churn'].value_counts())
✓ 0.0s

class distribution before augmentation:
churn
0    7963
1    2037
```

- In a real life also we only care about the people who are quitting or leaving (Exited) the bank, and we only want to analyze the patterns of those people.

## 2.3. Data Augmentation

I used data augmentation techniques to balance the dataset between the "Exited" and "Not Exited" classes. Since the original dataset had an imbalance between the number of customers who left the bank (Exited) and those who stayed (Not Exited), augmentation was applied to ensure the model receives equal representation from both classes. This step helps improve the model's ability to accurately predict customer churn without being biased toward the majority class.

```


58 # Check the class distribution after augmentation
59 print("class distribution after augmentation: ")
60 print(df_resampled['churn'].value_counts())
✓ 0.3s

```

```

class distribution after augmentation:
churn
1    7963
0    7963

```



```

Data Shape: (15926, 11)
Total Outliers in credit_score: 34 -- 0.21%
Total Outliers in age: 253 -- 1.59%
Total Outliers in tenure: 0 -- 0.0%
Total Outliers in balance: 0 -- 0.0%
Total Outliers in products_number: 64 -- 0.41%
Total Outliers in credit_card: 0 -- 0.0%
Total Outliers in active_member: 0 -- 0.0%
Total Outliers in estimated_salary: 0 -- 0.0%
Total Outliers in churn: 0 -- 0.0%
Cleaned Data Shape: (15575, 11)

```

The cleaned data shape was (15575, 11), indicating that outliers were addressed effectively to improve model performance.

## 2.4. Standardizing Data

Before training the Gradient Boosting model, the data was standardized to ensure that all features contributed equally to the model.

### Why Standardization?

Standardizing the data ensures that all features are on the same scale, preventing models from assigning disproportionate importance to features with larger ranges.

# 3. PyCaret for Algorithm Selection

## 3.1. Overview of PyCaret

PyCaret is an open-source, low-code machine learning library that streamline the process of finding the most suitable algorithm, PyCaret was used. PyCaret simplifies the comparison of machine learning models and automatically tunes hyperparameters.

## 3.2. Rationale for Using PyCaret

### Why PyCaret?

The key reason for using PyCaret in this workflow was to speed up the model comparison process. Rather than manually coding each model, PyCaret allowed us to automatically test a wide range of classifiers. This significantly reduces the time required to identify the best performing model for our churn prediction task.

```
save_model(tuned_model, 'best_churn_model')
```

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC	TT (Sec)
<b>gbc</b>	Gradient Boosting Classifier	0.8678	0.8684	0.4789	0.7797	0.5928	0.5193	0.5416	
<b>rf</b>	Random Forest Classifier	0.8609	0.8550	0.4455	0.7663	0.5629	0.4871	0.5127	
<b>lightgbm</b>	Light Gradient Boosting Machine	0.8608	0.8586	0.4878	0.7307	0.5842	0.5047	0.5200	
<b>et</b>	Extra Trees Classifier	0.8578	0.8481	0.4492	0.7428	0.5592	0.4806	0.5026	
<b>ada</b>	Ada Boost Classifier	0.8575	0.8468	0.4818	0.7177	0.5761	0.4945	0.5089	
<b>xgboost</b>	Extreme Gradient Boosting	0.8543	0.8509	0.4877	0.6972	0.5735	0.4891	0.5006	
<b>lr</b>	Logistic Regression	0.8314	0.7965	0.3284	0.6636	0.4388	0.3529	0.3834	
<b>ridge</b>	Ridge Classifier	0.8287	0.7954	0.2276	0.7399	0.3470	0.2797	0.3448	
<b>lda</b>	Linear Discriminant Analysis	0.8287	0.7954	0.3410	0.6409	0.4447	0.3545	0.3794	
<b>svm</b>	SVM - Linear Kernel	0.8205	0.7699	0.3143	0.6247	0.4056	0.3157	0.3468	
<b>knn</b>	K Neighbors Classifier	0.8194	0.7545	0.3314	0.5929	0.4246	0.3276	0.3474	
<b>nb</b>	Naive Bayes	0.8069	0.7870	0.4114	0.5268	0.4613	0.3460	0.3503	
<b>dummy</b>	Dummy Classifier	0.7987	0.5000	0.0000	0.0000	0.0000	0.0000	0.0000	
<b>dt</b>	Decision Tree Classifier	0.7963	0.6906	0.5137	0.4940	0.5032	0.3753	0.3756	
<b>qda</b>	Quadratic Discriminant Analysis	0.5223	0.5290	0.5074	0.1982	0.2536	0.0199	0.0270	



### 3.3. Algorithm Comparison

The table from PyCaret shows the performance of various models based on several metrics, including **Accuracy**, **AUC**, **Recall**, **Precision**, **F1-Score**, **Kappa**, and **MCC**. Here are the highlights:

- **Gradient Boosting Classifier (GBC)** performed the best with:
  - **Accuracy:** 0.8678
  - **AUC:** 0.8834
  - **Recall:** 0.4789
  - **Precision:** 0.7797
  - **F1-Score:** 0.5928
  - **Kappa:** 0.5198
  - **MCC:** 0.5416
- **Random Forest Classifier (RF)** and **LightGBM** also performed well but with slightly lower metrics compared to GBC.
- Other models like **Logistic Regression (LR)** and **Support Vector Machine (SVM)** showed moderate performance but were outperformed by tree-based models in terms of accuracy and AUC.
- **Decision Tree Classifier (DT)** had the lowest performance among the models tested with an accuracy of **0.7963**.

These results indicate that **Gradient Boosting Classifier (GBC)** was the best choice for this churn prediction task. The decision to proceed with GBC was driven by its higher accuracy and AUC scores, which are critical for predicting customer churn.

# 4. Modelling

## 4.1. Gradient Boosting Classifier

Gradient Boosting Classifier was used to build a stronger model by combining several weak learners (decision trees). This ensemble method reduces variance and improves model accuracy.

- **Optimized Accuracy:** 84%
- **AUC-ROC:** 0.9197
- **Best Hyperparameters:**
  - Learning Rate: 0.15
  - Max Depth: 4
  - Min Samples Leaf: 2
  - Min Samples Split: 5
  - N Estimators: 250

Classification Report:

```
Optimized Accuracy: 0.8371495827091804
Classification Report:

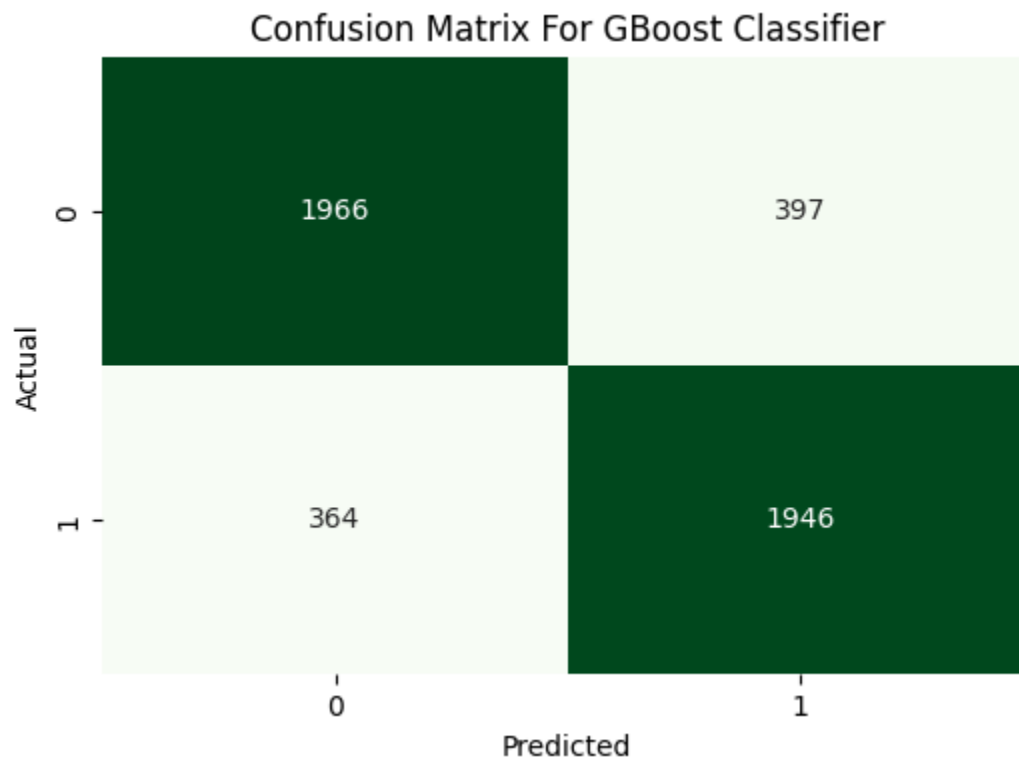
```

	precision	recall	f1-score	support
0	0.84	0.83	0.84	2363
1	0.83	0.84	0.84	2310
accuracy			0.84	4673
macro avg	0.84	0.84	0.84	4673
weighted avg	0.84	0.84	0.84	4673

```
AUC-ROC: 0.9196970612967228
```

## Confusion Matrix for GBC

- True Positives (TP): 1946
- False Positives (FP): 397
- True Negatives (TN): 1966
- False Negatives (FN): 364



This confusion matrix illustrates the model's performance in classifying the churn variable.

## 4.2. Logistic Regression

Logistic Regression was chosen as the baseline model due to its simplicity and interpretability. It provides a quick way to get insights into the dataset, especially for binary classification tasks.

- **Accuracy: 82%**
- **AUC-ROC: 0.9026**

### Classification Report:

```
Accuracy: 0.8195826645264848
Classification Report:
      precision    recall  f1-score   support


     0       0.82      0.82      0.82     1532
     1       0.82      0.82      0.82     1583

 accuracy          0.82          0.82          0.82          3115
  macro avg       0.82      0.82      0.82          3115
 weighted avg     0.82      0.82      0.82          3115

AUC-ROC: 0.9026483244789201
```

## Confusion Matrix for Logistic Regression

- True Positives (TP): 1303
- False Positives (FP): 282
- True Negatives (TN): 1250
- False Negatives (FN): 280



This confusion matrix illustrates the model's performance in classifying the churn variable.

## 4.3. Support Vector Classifier (SVC)

The SVC was chosen for its ability to handle high-dimensional data and separate classes that are not linearly separable.

- **Accuracy:** 81.86%
- **AUC-ROC:** 0.9027

Classification Report:

```
Accuracy: 0.8186195826645265
Classification Report:
              precision    recall  f1-score   support


     0           0.82         0.82         0.82         1532
     1           0.82         0.82         0.82         1583

 accuracy                   0.82         3115
 macro avg           0.82         0.82         0.82         3115
 weighted avg        0.82         0.82         0.82         3115

AUC-ROC: 0.9027039910009912
```

## Confusion Matrix for SVC

- True Positives (TP): 1301
- False Positives (FP): 283
- True Negatives (TN): 1249
- False Negatives (FN): 282

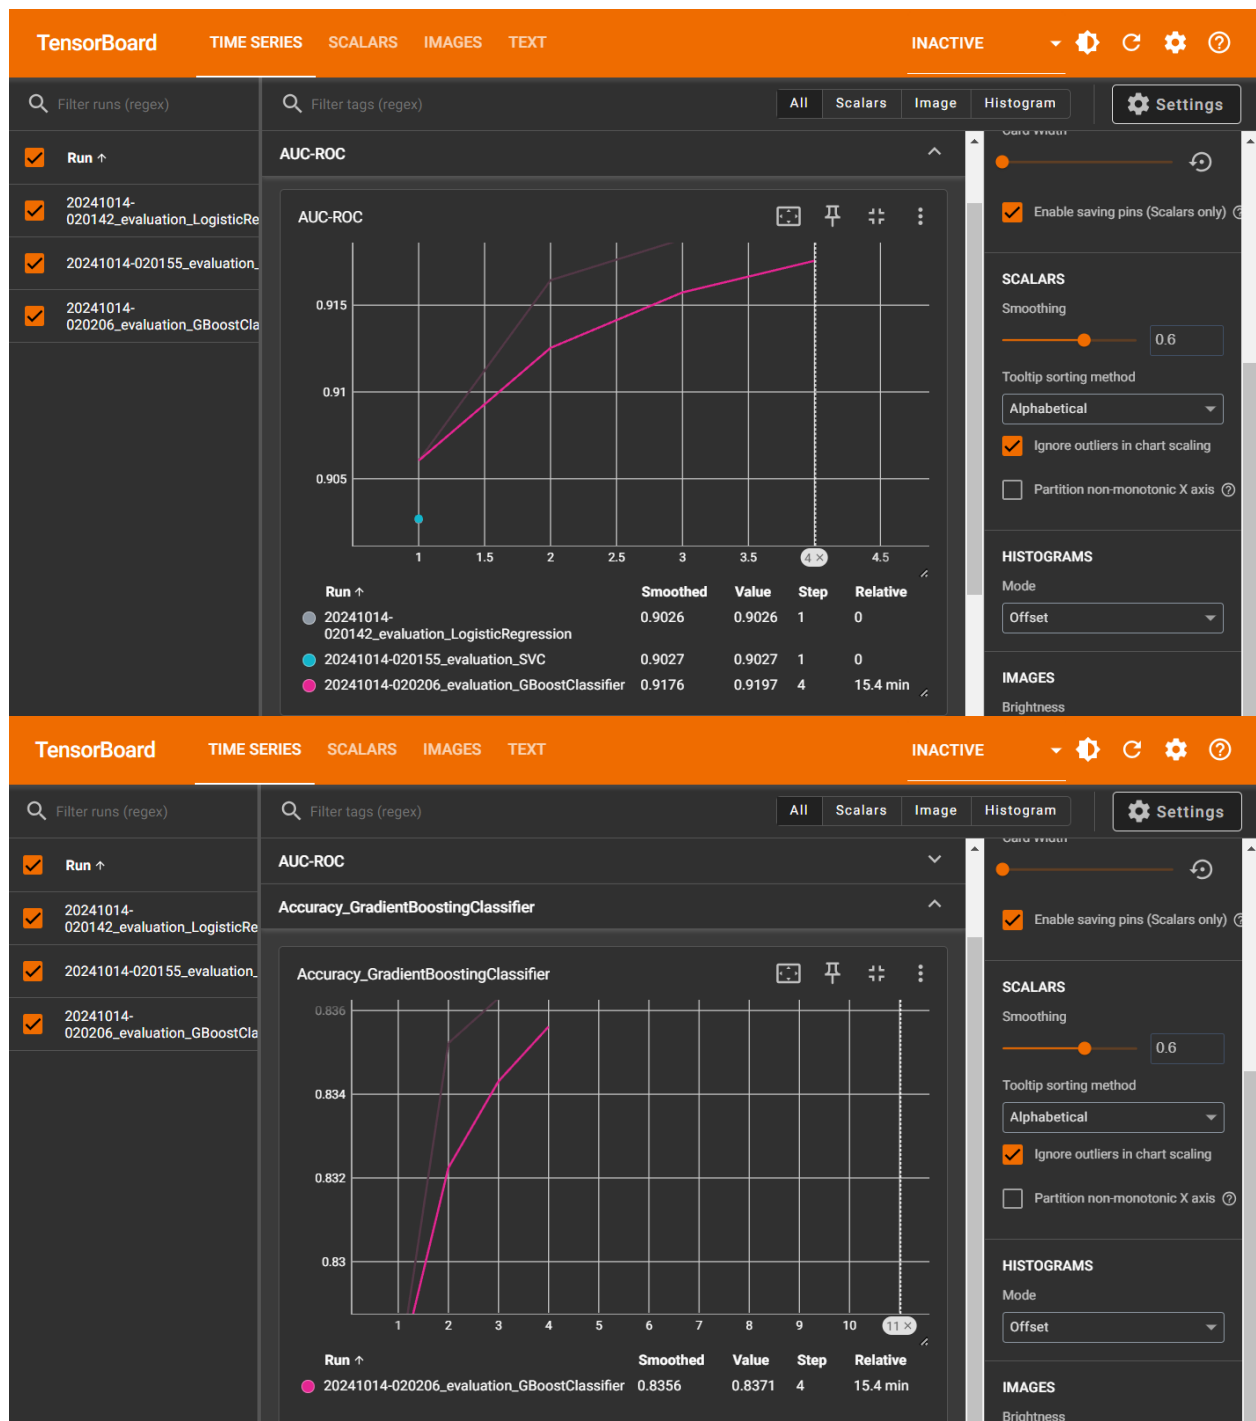


This confusion matrix illustrates the model's performance in classifying the churn variable.

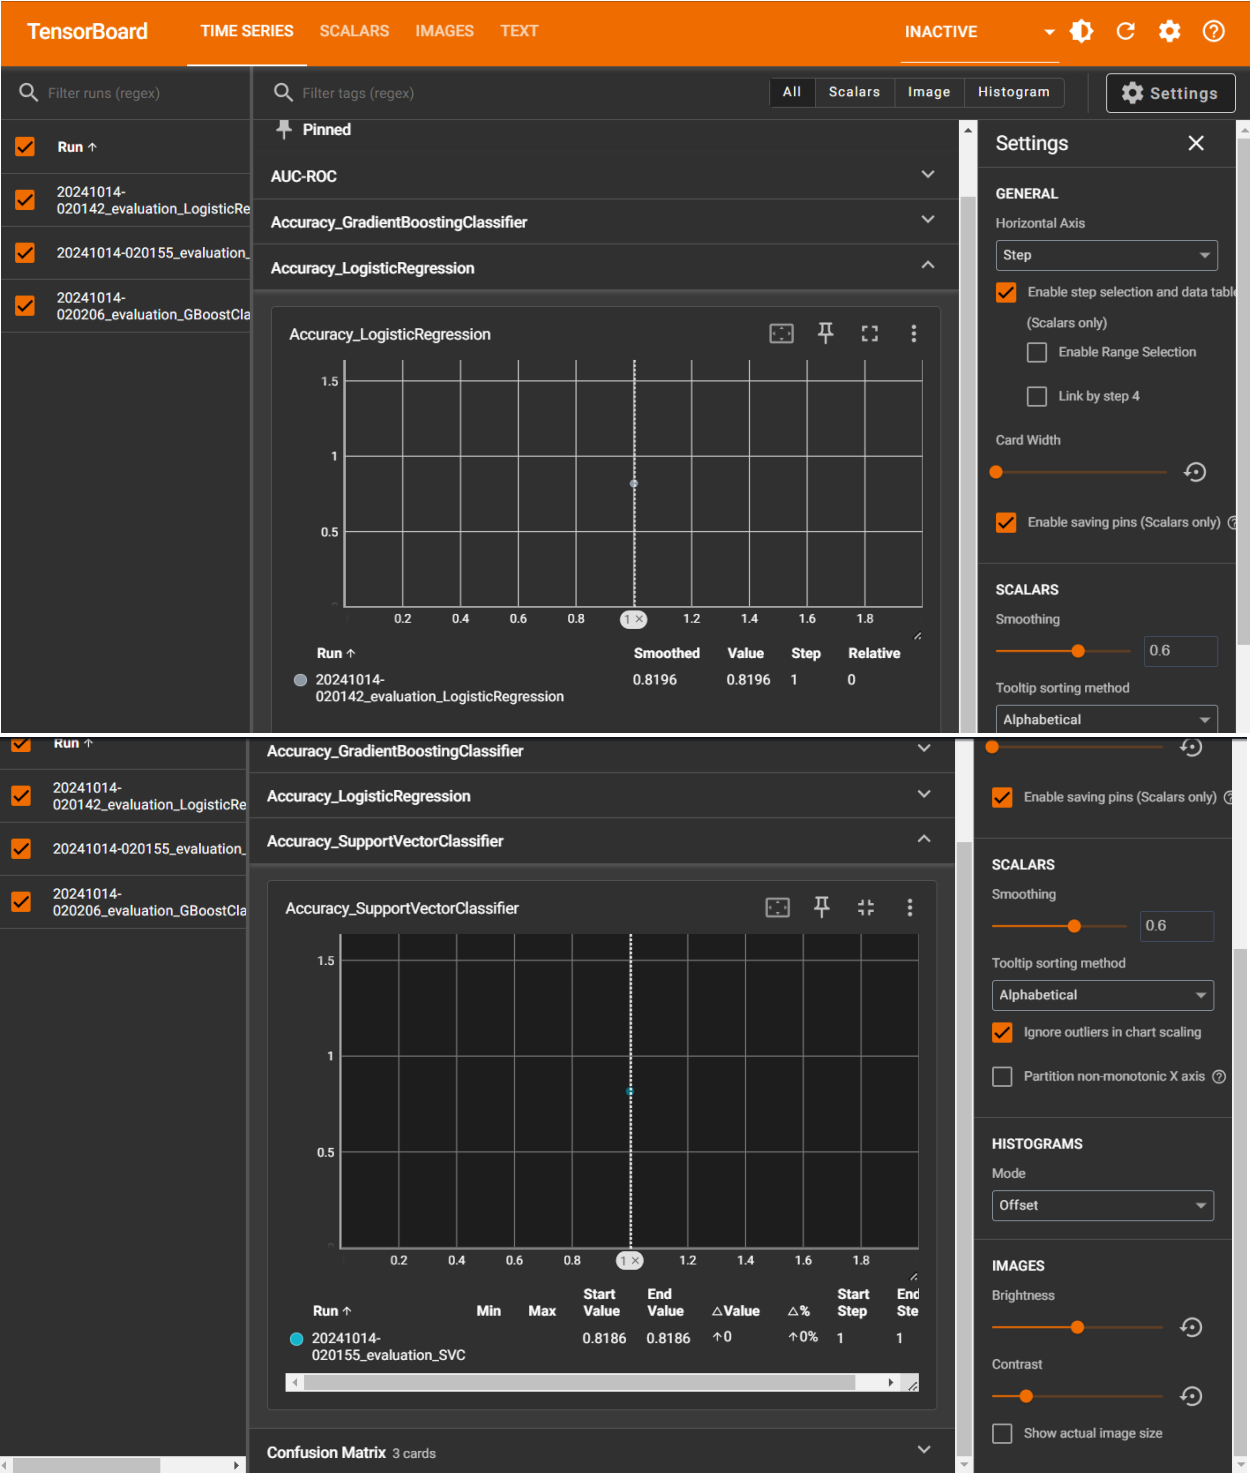
## 5. TensorBoard for Workflow Visualization

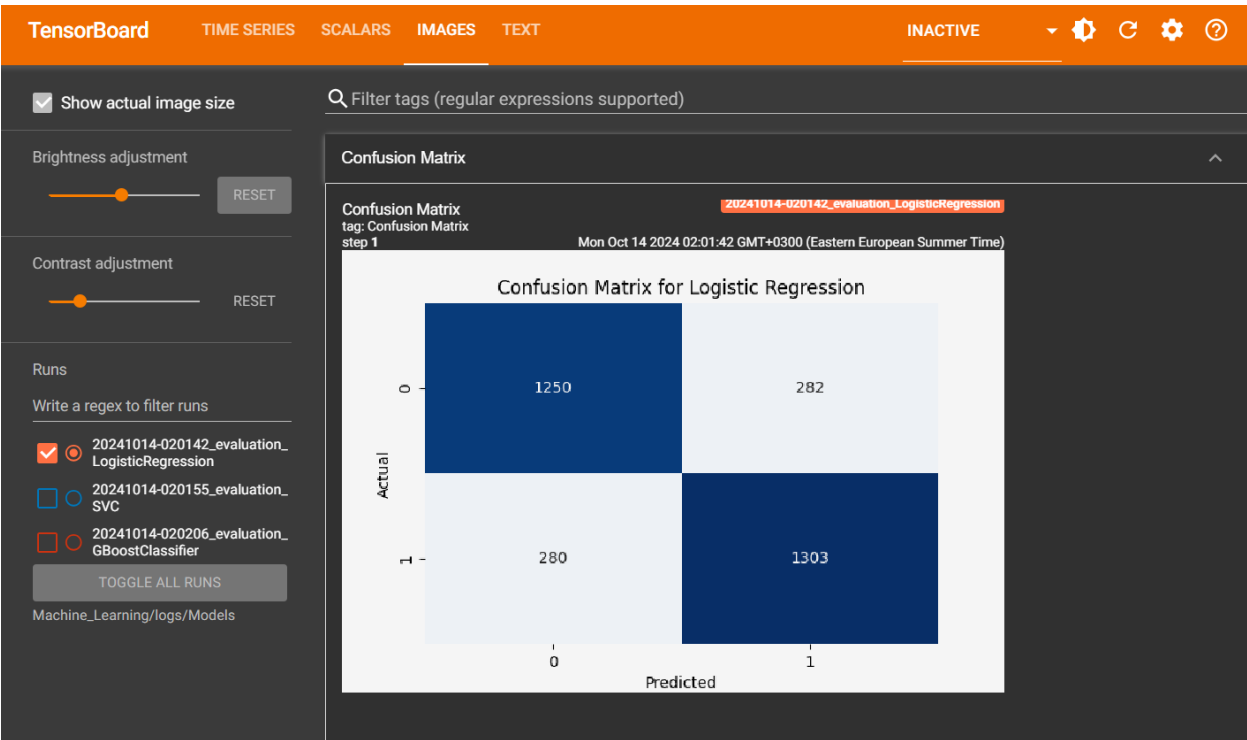
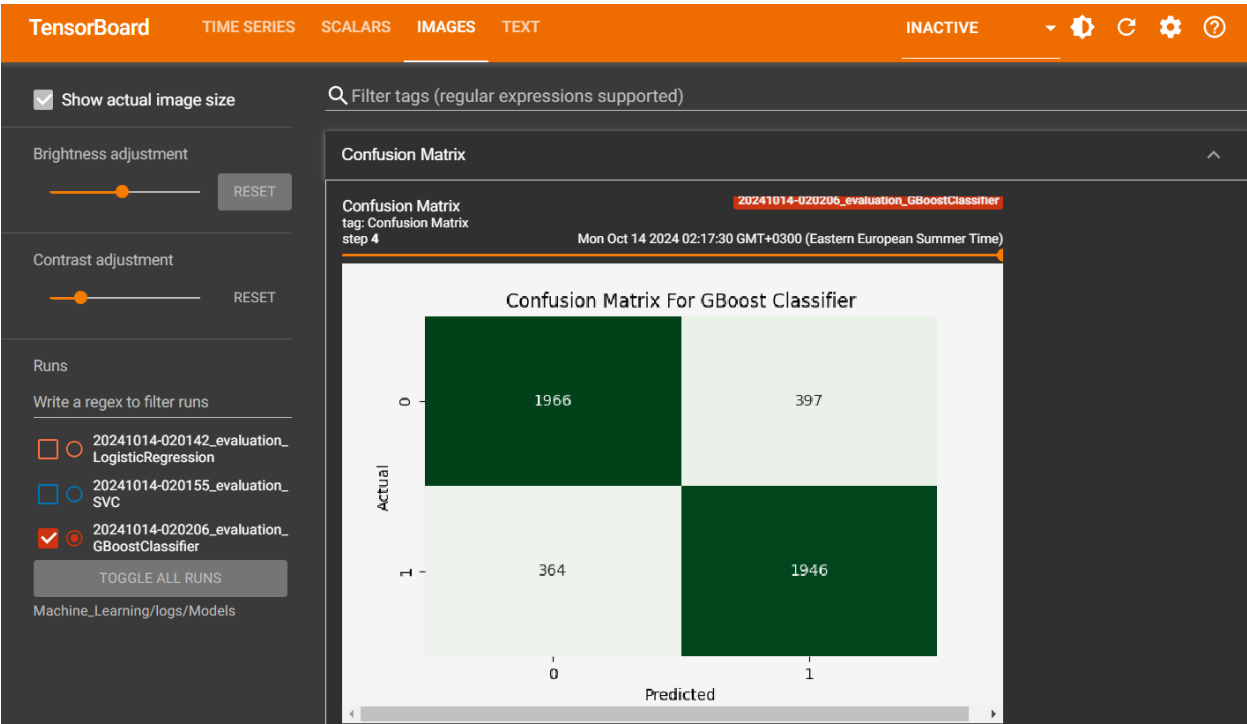
TensorBoard was used to visualize the model training process, including metrics such as accuracy, AUC-ROC scores, and confusion matrices for each model. It helped in monitoring the workflow and provided real-time insights into model performance.

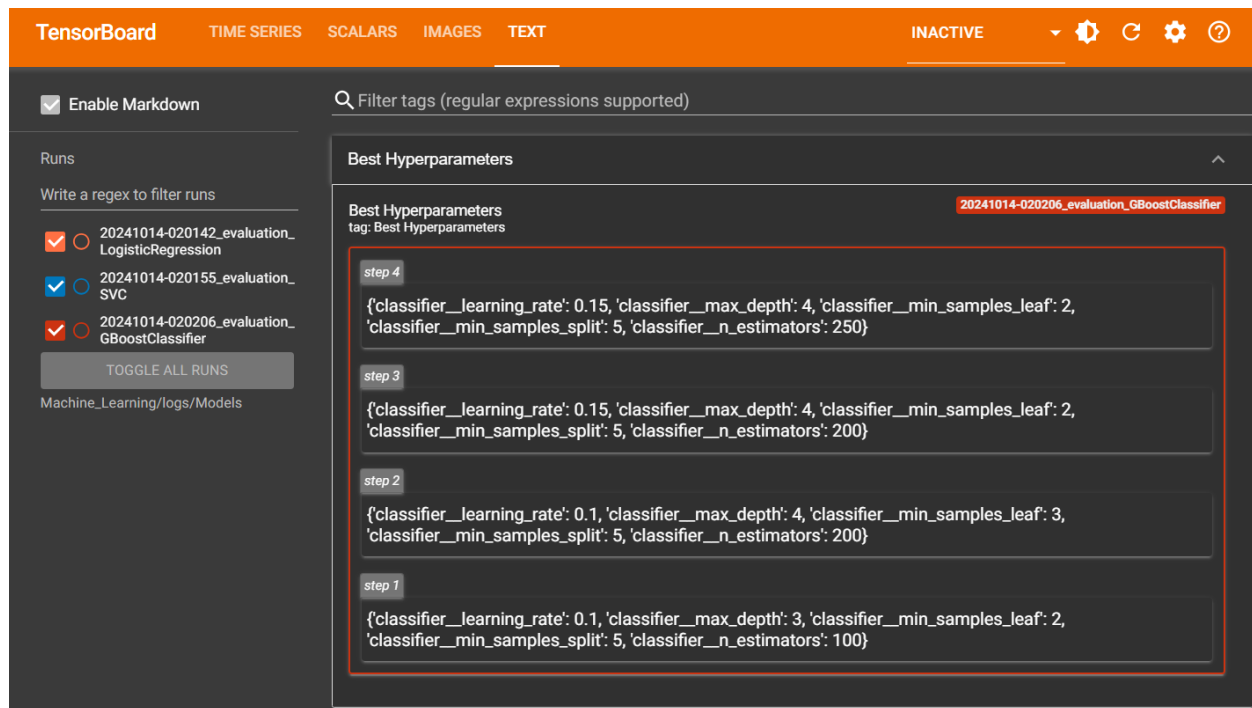
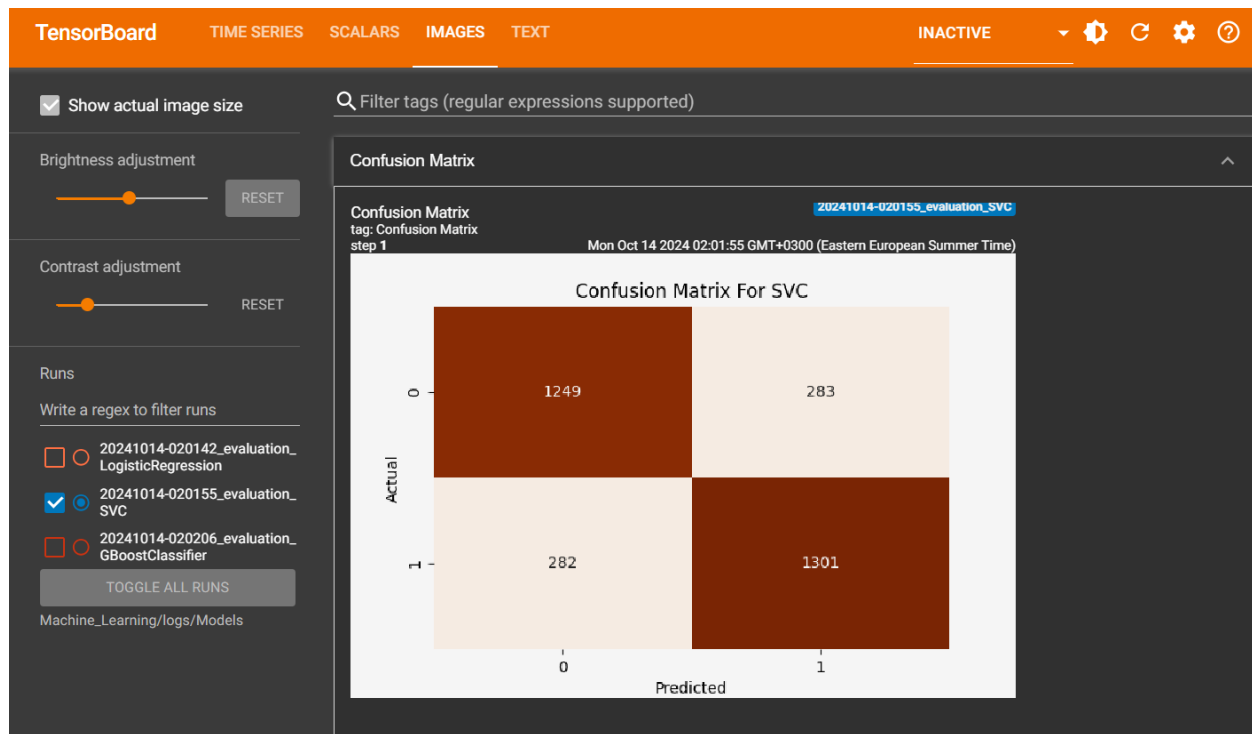
TensorBoard Logs Created At: log\_dir











## 6. Conclusion

This report covered the workflow for building and evaluating machine learning models using Gradient Boosting Classifier, Logistic Regression, and SVC. Each model provided valuable insights, with Gradient Boosting emerging as the most accurate model. TensorBoard facilitated real-time monitoring of the workflow, contributing to better decision-making during model development. Future improvements could include exploring more advanced models and fine-tuning the existing ones to enhance accuracy further.

### Summary of Findings

Throughout the process of building a machine learning model to predict customer churn, several key steps were undertaken, including data preparation, feature engineering, and the use of **PyCaret** for algorithm comparison. The analysis provided important insights into the performance of various models, particularly in terms of their ability to accurately predict which customers are likely to leave the bank.

- **Data Preparation and Augmentation:** The dataset was balanced using data augmentation techniques, ensuring fair representation between the "Exited" and "Not Exited" classes.
- **Handling Outliers:** Outliers were detected and handled in key variables such as **credit\_score** and **age**, improving the robustness of the model.
- **Algorithm Comparison:** Using PyCaret, multiple machine learning algorithms were tested, including **Gradient Boosting Classifier (GBC)**, **Random Forest (RF)**, **Logistic Regression (LR)**, and **Support Vector Classifier (SVC)**. Among them, **Gradient Boosting Classifier** outperformed all others with the highest accuracy and AUC score.
- **Model Metrics:** The models were evaluated based on various metrics such as **Accuracy**, **AUC**, **Recall**, **Precision**, and **F1-Score**. The GBC model demonstrated the best balance between precision and recall, making it the most suitable model for this churn prediction task.